




Chapter 8: Deadlocks

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock
- Combined Approach to Deadlock Handling





The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Example
 - ◆ System has 2 tape drives.
 - ◆ P_1 and P_2 each hold one tape drive and each needs another one.
- Example
 - ◆ semaphores A and B , initialized to 1


P_0
 $wait(A);$
 $wait(B);$

P_1
 $wait(B)$
 $wait(A)$





Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation(every thin stops) is possible.





System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - ◆ request
 - ◆ use
 - ◆ release






Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_{n-1}\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_0 is waiting for a resource that is held by P_0 .





Resource-Allocation Graph

A set of vertices V and a set of edges E .


■ V is partitioned into two types:

◆ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.

◆ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.

■ request edge – directed edge $P_i \rightarrow R_j$

■ assignment edge – directed edge $R_j \rightarrow P_i$





Resource-Allocation Graph (Cont.)


■ Process





■ Resource Type with 4 instances



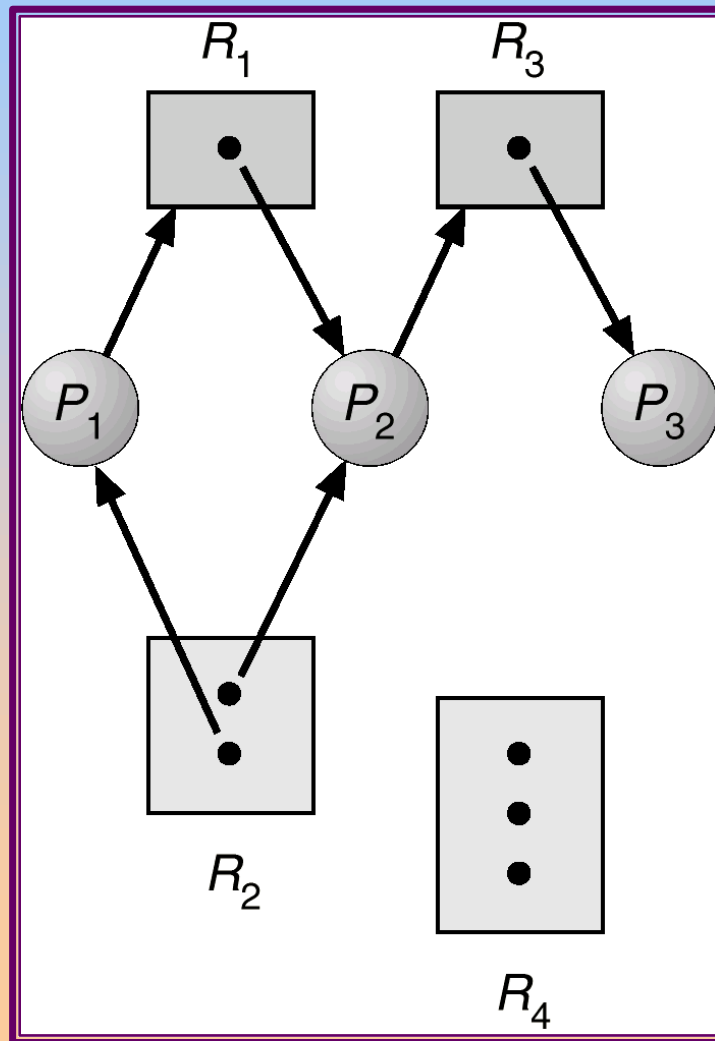
■ P_i requests instance of R_j



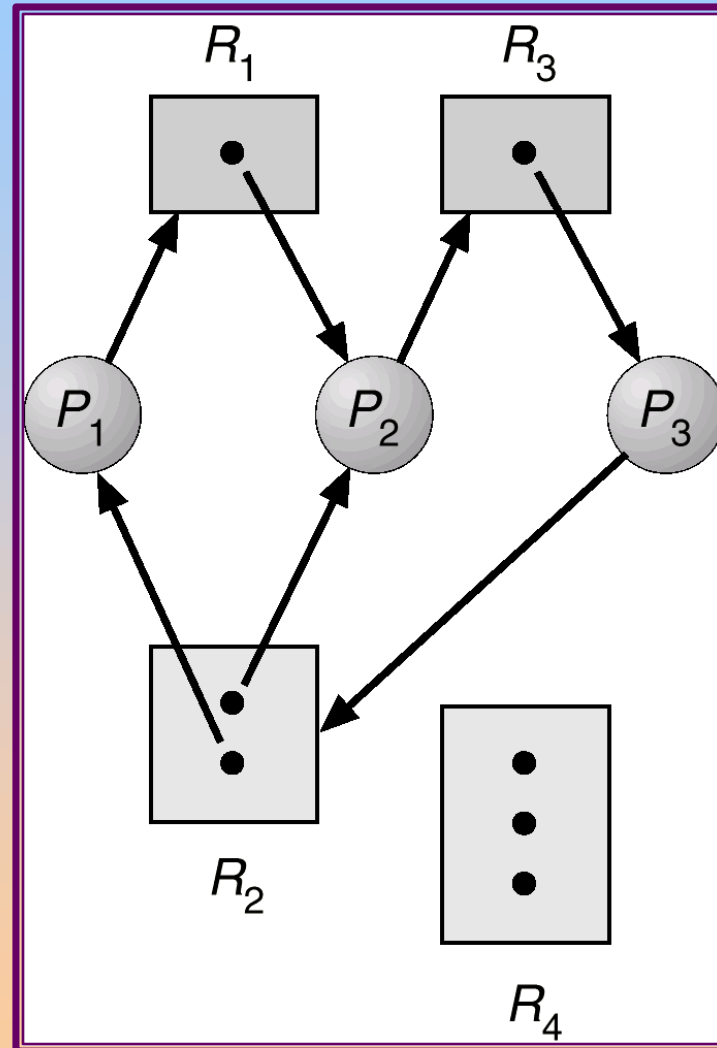
■ P_i is holding an instance of R_j




Example of a Resource Allocation Graph



Resource Allocation Graph With A Deadlock




Resource Allocation Graph With A Cycle But No Deadlock





Basic Facts


- If graph contains no cycles → no deadlock.
- If graph contains a cycle →
 - ◆ if only one instance per resource type, then deadlock.
 - ◆ if several instances per resource type, possibility of deadlock.
- ◆ **Definition of: instance** (1) A single copy of a running program. Multiple **instances** of a program **mean** that the program has been loaded into memory several times.





Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
- Allow the system to enter a deadlock state and then recover.
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX.





Deadlock Prevention

Restrain the ways request can be made.

- **Mutual Exclusion** – not required for sharable resources;
The process must hold for nonsharable resources.
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources.
 - ◆ Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.
 - ◆ Low resource utilization; starvation possible.






Deadlock Prevention (Cont.)

■ No Preemption –

- ◆ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- ◆ Preempted resources are added to the list of resources for which the process is waiting.
- ◆ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

■ Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.





Deadlock Avoidance

Requires that the system has some additional *a priori* information available.


- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.





Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a *safe state*.
- System is in safe state if there exists a safe sequence of all processes.
- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$.
 - ◆ If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
 - ◆ When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - ◆ When P_i terminates, P_{i+1} can obtain its needed resources, and so on.






Basic Facts

- If a system is in safe state → no deadlocks.
- If a system is in unsafe state → possibility of deadlock.
- Avoidance → ensure that a system will never enter an unsafe state.





Safe, Unsafe , Deadlock State






Resource-Allocation Graph Algorithm


- *Claim edge* $P_i \rightarrow R_j$ indicated that process P_j may request resource R_j ; represented by a dashed line.
- Claim edge converts to request edge when a process requests a resource.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.



Resource-Allocation Graph For Deadlock Avoidance




Unsafe State In Resource-Allocation Graph





Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme





Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
 - ◆ Priority of the process.
 - ◆ How long process has computed, and how much longer to completion.
 - ◆ Resources the process has used.
 - ◆ Resources process needs to complete.
 - ◆ How many processes will need to be terminated.
 - ◆ Is process interactive or batch?






Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.





Combined Approach to Deadlock Handling

- Combine the three basic approaches

- ◆ prevention
- ◆ avoidance
- ◆ detection

allowing the use of the optimal approach for each of resources in the system.

- Partition resources into hierarchically ordered classes.

- Use most appropriate technique for handling deadlocks within each class.

